

[Get started](#)[Open in app](#)[Follow](#)

572K Followers



This is your **last** free member-only story this month. [Sign up for Medium and get an extra one](#)

Sentiment Analyzer with BERT (build, tune, deploy)

Brief description of how I developed sentiment analyzer. It covers text preprocessing, model building, tuning, API, frontend creation and containerization.




Zuzanna Deutschman · Jul 24, 2020 · 4 min read ★

Sentiment Analyzer

Provide sentiment score for your text.

Please enter the text



Dataset

I used the [dataset published by The Stanford NLP Group](#). I merged two files, namely 'dictionary.txt' including 239,232 text fragments and 'sentiment_labels.txt' containing the sentiment scores assigned to the various text fragments.

Text preprocessing with regular expressions

To clean the text, I usually use a bunch of functions containing regular expressions. In `common.py` you can find all of them, for example `remove_nonwords` described below:

```
1  def remove_nonwords(df, column):
2      """
3      Replace non-words from the beginning and end of the string
4      Parameters
5      -----
6      df : data frame
7      column : column name with input text
8
9      Returns
10     -----
11     df : data frame with cleaned text
12     """
13     reg = r"^[^a-zA-Z]*|[^a-zA-Z]*$"
14     df[column] = df[column].str.replace(reg, "")
15     return df
```

`remove_nonwords.py` hosted with ❤ by GitHub

[view raw](#)

Similar functions were used for empty rows, special signs, numbers and html code removal.

After text cleaning, it's time for BERT embeddings creation. For that purpose, I used [bert-as-service](#). It is very simple and consists of only 3 steps: download a pre-trained model, start the BERT service and use client for sentence encodings of specified length.

```
1 def BERT_embeddings(train_x, test_x, valid_x, valid_y):
2     bc = BertClient(ip='ip')
3     train_seq_x = bc.encode(list(train_x.values))
4     test_seq_x = bc.encode(list(test_x.values))
5     valid_seq_x = bc.encode(list(valid_x.values))
6
7     train_x = pd.DataFrame(data=train_seq_x, index=train_x.index)
8     test_x = pd.DataFrame(data=test_seq_x, index=test_x.index)
9     valid_x = pd.DataFrame(data=valid_seq_x, index=valid_x.index)
10    return train_x, test_x, valid_x
```

BERT_embeddings.py hosted with ❤ by GitHub

[view raw](#)

There are multiple parameters that can be setup, when running a service. For example, to define `max_seq_len`, I calculated 0.9 quantile of train data length.

```
1 def max_sequence_length(train_x):
2     return int(train_x.str.split().str.len().quantile(0.9))
```

max_seq_len.py hosted with ❤ by GitHub

[view raw](#)

Preprocessed data has a form of data frame containing 768 features. For full code, please go to `nlp_preprocess.py`.

Model building with Keras

In this part, we build and train the model on different parameters. Let's assume we want 5-layers neural network as below. We will parametrize `batch_size`, number of epochs, number of nodes in the first 4 dense layers and 5 dropout layers.

```
1 def build_model(train_x, batch_size=128,
2                 dense_1_nodes=256, dense_2_nodes=256, dense_3_nodes=256, dense_4_nodes=256,
3                 dropout_1_size=0.1, dropout_2_size=0.1, dropout_3_size=0.1, dropout_4_size=0.1,
```

```
4 model = Sequential()
5
6 # The Input Layer :
7 model.add(Dense(dense_1_nodes, kernel_initializer='normal',input_dim = train_x.shape[1], act
8 model.add(Dropout(rate=dropout_1_size))
9
10 # The Hidden Layers :
11 model.add(Dense(dense_2_nodes, kernel_initializer='normal',activation='relu'))
12 model.add(Dropout(dropout_2_size))
13
14 model.add(Dense(dense_3_nodes, kernel_initializer='normal',activation='relu'))
15 model.add(Dropout(dropout_3_size))
16
17 model.add(Dense(dense_4_nodes, kernel_initializer='normal',activation='relu'))
18 model.add(Dropout(dropout_4_size))
19
20 # The Output Layer :
21 model.add(Dense(units = 1, activation="sigmoid"))
22 model.add(Dropout(dropout_5_size))
23 model.compile(loss='mean_squared_error', optimizer='rmsprop', metrics = ['mse', 'mae'])
24 print(model.summary)
25 return model
```

build_model.py hosted with ❤ by GitHub

[view raw](#)

```
1 def model_training(train_x, train_y, valid_x, valid_y, epochs, batch_size,
2                     dense_1_nodes, dense_2_nodes, dense_3_nodes, dense_4_nodes,
3                     dropout_1_size, dropout_2_size, dropout_3_size, dropout_4_size, dropout_5_size):
4     model = build_model(train_x, batch_size,
5                         dense_1_nodes, dense_2_nodes, dense_3_nodes, dense_4_nodes,
6                         dropout_1_size, dropout_2_size, dropout_3_size, dropout_4_size, dropout_5_size)
7     history = model.fit(train_x.values, train_y.values, epochs=epochs, batch_size=batch_size, ver
8     model.save('model.h5')
9     return model, history
```

model_training.py hosted with ❤ by GitHub

[view raw](#)

Model tuning with Sacred

Now we can tune the parameters. We will use `sacred` module. Key points here are:

1. Create an Experiment and add Observer

First we need to create an experiment and observer that logs all kinds of information. It's very simple!

```
1 from sacred import Experiment
2 from sacred.observers import MongoObserver
3
4 ex = Experiment()
5 ex.observers.append(MongoObserver(
6     url='mongodb://192.168.1.2:27017'))
```

create_experiment.py hosted with ❤ by GitHub

[view raw](#)

2. Define the main function

The `@ex.automain` decorator defines and runs the main function of the experiment when we run the Python script.

```
1 @ex.automain
2 def main(epochs, batch_size,
3         dense_1_nodes, dense_2_nodes, dense_3_nodes, dense_4_nodes,
4         dropout_1_size, dropout_2_size, dropout_3_size, dropout_4_size, dropout_5_size):
5
6     train_x, train_y, test_x, test_y, valid_x, valid_y = load_model()
7     model, history = model_training(train_x, train_y, valid_x, valid_y, epochs=epochs, batch_size=batch_size,
8                                     dense_1_nodes=dense_1_nodes, dense_2_nodes=dense_2_nodes, dense_3_nodes=dense_3_nodes, dense_4_nodes=dense_4_nodes,
9                                     dropout_1_size=dropout_1_size, dropout_2_size=dropout_2_size, dropout_3_size=dropout_3_size, dropout_4_size=dropout_4_size, dropout_5_size=dropout_5_size)
10    model_testing(test_x, test_y, model, history)
```

main.py hosted with ❤ by GitHub

[view raw](#)

3. Add the Configuration parameters

We will define them through Config Scope.

```
1 @ex.config
2 def cfg():
```

```
3     epochs=30
4     batch_size=128
5     dense_1_nodes=256
6     dense_2_nodes=256
7     dense_3_nodes=256
8     dense_4_nodes=256
9     dropout_1_size=0.1
10    dropout_2_size=0.1
11    dropout_3_size=0.1
12    dropout_4_size=0.1
13    dropout_5_size=0.3
```

cfg.py hosted with ❤ by GitHub

[view raw](#)

4. Add metrics

In our case here I want to know the **MAE** and **MSE**. We can use the Metrics API for that.

```
1  def model_testing(test_x, test_y, model, history):
2      pred_y = model.predict(test_x)
3      np.savetxt("data/pred_y.csv", pred_y, delimiter=",")
4      mae = mean_absolute_error(pred_y, test_y)
5      mse = mean_squared_error(pred_y, test_y)
6      ex.log_scalar("MAE", mae)
7      ex.log_scalar("MSE", mse)
```

model_testing.py hosted with ❤ by GitHub

[view raw](#)

5. Run the experiment

Functions from the previous steps are stored in `model_experiment.py` script. In order to run our experiment for bunch of parameters, we create and run `run_sacred.py`. For all possible permutations, MAE and MSE will be saved in MongoDB.

```
1  from model_experiment import ex
2  import itertools
3
4  batch_size_values =[64, 128]
5  dense_1_nodes =[64, 128]
6  dense_2_nodes =[64, 128]
7  dense_3_nodes =[64, 128]
8  dense_4_nodes =[64, 128]
```

```
9 dropout_1_size = [0.1, 0.3]
10 dropout_2_size = [0.1, 0.3]
11 dropout_3_size = [0.1, 0.3]
12 dropout_4_size = [0.1, 0.3]
13 dropout_5_size = [0.1, 0.3]
14 epochs_values = [30]
15
16
17 for epochs, batch_size, dense_1_nodes, dense_2_nodes, dense_3_nodes, dense_4_nodes, dropout_1_si
18     in itertools.product(epochs_values, batch_size_values, dense_1_nodes, dense_2_nodes, dense_3
19     ex.run(config_updates={'dropout_5_size': dropout_5_size, 'batch_size': batch_size, 'epochs':
20         'dense_1_nodes': dense_1_nodes, 'dense_2_nodes': dense_2_nodes,
21         'dense_3_nodes': dense_3_nodes, 'dense_4_nodes': dense_4_nodes,
22         'dropout_1_size': dropout_1_size, 'dropout_2_size': dropout_2_size,
23         'dropout_3_size': dropout_3_size, 'dropout_4_size': dropout_4_size
24         })
```

run_sacred.py hosted with ❤ by GitHub

[view raw](#)

The best result I got is 9% of MAE score. That means that our sentiment analyzer works pretty good. We can check it with `model_inference` function.

```
1 def model_inference(model, sentence):
2     bc = BertClient()
3     sentence = bc.encode(sentence)
4     score = model.predict(sentence)
5     pred_y = np.loadtxt(os.path.normpath("data/pred_y.csv"), delimiter=",")
6     min_score, max_score = min_max_sentiment(pred_y)
7     norm_score = (score - min_score)/(max_score - min_score)
8     return norm_score
```

model_inference.py hosted with ❤ by GitHub

[view raw](#)

Please note that the score is normalized so that outlier values can be also obtained. After model is saved, we can build a Web API!

Web API creation with Flask

Now we want to create an API that runs the code in the function and displays the returned result in the browser.

```
1 @app.route('/score', methods = ['PUT'])
2 @cross_origin()
3 def score():
4     print(request)
5     text = request.json['text']
6     df = pd.DataFrame(columns=['text']).append({'text': text}, ignore_index=True)
7     df = text_cleaning(df, 'text')
8     preproc_text = df.iloc[0][0]
9     score = float(model_inference(model, [preproc_text]))
10    return {"score": score}
```

app.py hosted with ❤ by GitHub

[view raw](#)

The syntax `@app.route('/score', methods=['PUT'])` lets Flask know that the function, `score`, should be mapped to the *endpoint/score*. The `methods` list is a keyword argument that tells us what kind of HTTP requests are allowed. We'll be using `PUT` requests to receive sentences from a user. In function `score`, we get a score in dictionary form, since it can be easily converted to a JSON string. Full code is available in `api.py`.

Frontend

For web interface, three files were created:

- `index.html` provides the basic structure of the site: title, description, input text area and circle with score.
- `style.css` is used to style the website.
- `index.js` provides interactivity. It is responsible for reading user input, handling API requests and presenting calculated score. Three main functions here are:

```
1 function textChangedHandler(text) {
2     console.log(text)
3     fetch("http://127.0.0.1:5000/score", {
```



```
4     method: "PUT",
5     body: JSON.stringify({"text": text}),
6     headers: {
7       'Accept': 'application/json',
8       'Content-Type': 'application/json'},
9   }).then(parseResponse)
10 }
11
12 const textChanged = debounce(textChangedHandler, 250);
13
14 function parseResponse(response){
15   console.log(response)
16   response.json().then(handleJSON)
17 }
18
19 function handleJSON(response){
20   let score = (response.score * 100).toFixed(0)+'%'
21   console.log(score)
22   document.getElementById("score").textContent = score
23   document.getElementById("circle").style.backgroundColor = hsv2hex({h: response.score*120, s:
24 }
```

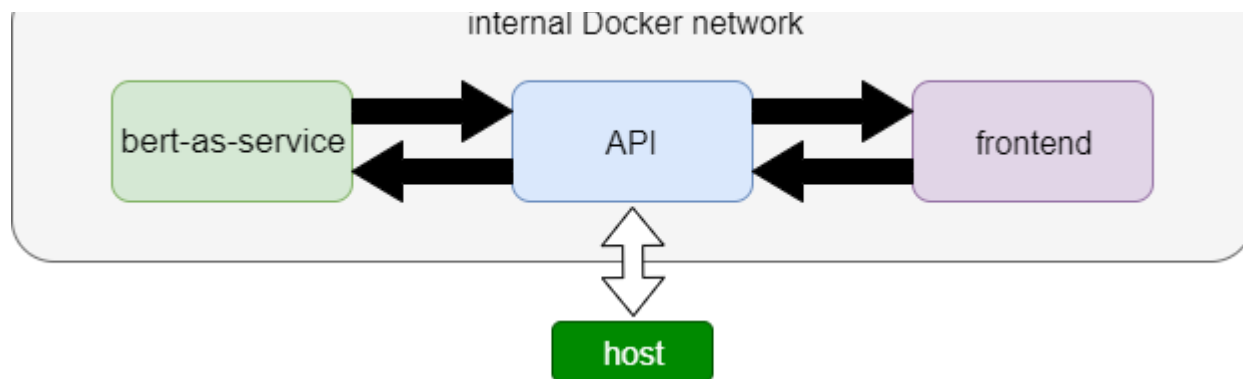
index.js hosted with ❤ by GitHub

[view raw](#)

For gradient **HSV** model was used. Saturation and Value are constants. Hue corresponds to score value. Changing hue in range [0;120] yields smooth colour change from red to yellow to green.

Docker containerization

The brilliance of Docker is that, once you package an application and all its dependencies into container, you ensure it will run in any environment. It is generally recommended to separate areas of concern by using one service per container. In my small app there are 3 parts that should be combined: bert-as-service, application and frontend. The tool that helps you build Docker images and run containers is **Docker Compose**.



Steps that we need to do to dockerize our code:

- Create separate folders for bert-as-service, api and frontend,
- Put there relevant files,
- Add `requirements.txt` and `Dockerfile` to each folder. The first file should cover all needed libraries that will be installed via command in the second file. Its format is described in docker [documentation](#)
- Create `docker-compose.yaml` in the 3 folders directory. Define the 3 services that make up the app in this file, so they can be run together in an isolated environment.

Now we are ready to build and run our application! Please see the sample outputs below.

Brave and funny movie with fantastic screenplay. I highly recommend it for the weekend.

84%

Just another Tom Hanks movie.

45%

The worst thing I've ever seen. Nightmare.

16%

As usual, please feel free to view the full code on my Gitlab.

Projects - Zuzanna / Sentiment Analysis with BERT

GitLab.com

gitlab.co

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

Your email

Get this newsletter

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

[Data Science](#) [Deep Learning](#) [NLP](#) [Deployment](#) [Machine Learning](#)

[About](#) [Help](#) [Legal](#)

Get the Medium app

